

```
1 import pandas as pd
2 import numpy as np #9
3 import matplotlib.pyplot as plt #9
4 from sklearn.model_selection import train_test_split #4
5 from sklearn.preprocessing import StandardScaler #5
6 from sklearn.linear_model import LinearRegression #6 & #9
7 from sklearn.metrics import mean_squared_error #6 & #7
8 from sklearn.tree import DecisionTreeRegressor #7
9 from sklearn.ensemble import RandomForestRegressor #8
10
11 # Ignore printing warnings for general readability
12 import warnings
13 warnings.filterwarnings('ignore')

1 #Making a list of missing value types
2 missing_values = ["n/a", "na", "--", "...", "NaN"]
3
4 # 1. # Read the Excel file using the specified engine and convert to CSV
5 data = pd.read_excel('1553768847_housing.xlsx', na_values=missing_values)
6 data.to_csv('1553768847_housing.csv', index=False)
7 data.head(20)
```



	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-122.23	37.88	41	880	129.0	322	126	8.3252
1	-122.22	37.86	21	7099	1106.0	2401	1138	8.3014
2	-122.24	37.85	52	1467	190.0	496	177	7.2574
3	-122.25	37.85	52	1274	235.0	558	219	5.6431
4	-122.25	37.85	52	1627	280.0	565	259	3.8462

```
1 # 2. Handle missing values by filling with column mean
```

```
2 data_filled = data.fillna(data.mean())
```

```
1 #Identify the data variables which are categorical
```

```
2 #Categorical(classified as strings or integers) columns can be further classified into:
```

```
3 #-nominal(no inherent order/ranking since no specific order)
```

```
4 #-ordinal variables (specific order/ranking) or satisfaction ratings (e.g., low < medium < high)
```

```
5
```

```
6 #Continuous any numerical value(float/int) within a specific range(eg:fractional or decimal values/ ie: age, height, weight, or temperature)
```

```
7 Categorical_Columns = []
```

```
8 Numerical_Columns = []
```

```
9 for column in data_filled.columns:
```

```
10     if len(data_filled[column].unique()) <= 8:
```

```
11         Categorical_Columns.append(column)
```

```
12     else:
```

```
13         Numerical_Columns.append(column)
```

```
14
```

```
15 print("\n Categorical Variables are : " , Categorical_Columns)
```

```
16 print(" \n Numerical Variables are : " , Numerical_Columns)
```

```
Categorical Variables are : ['ocean_proximity']
```

```
Numerical Variables are : ['longitude', 'latitude', 'housing_median_age', 'total_rooms', 'total_bedrooms', 'population', 'households', 'median_income',
```

```
1 # 3. Encode categorical data: Convert categorical columns to numerical data
```

```
2 data_encoded = pd.get_dummies(data_filled, columns=['ocean_proximity'])
```

```
3
```

```
4 # Display the first few rows of the encoded dataset
```

```
5 data_encoded.head()
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income
0	-122.23	37.88	41	880	129.0	322	126	8.3252
1	-122.22	37.86	21	7099	1106.0	2401	1138	8.3014
2	-122.24	37.85	52	1467	190.0	496	177	7.2574

```

1 # Split the dataset into features (X) and target variable (y)
2 X = data_encoded.drop('median_house_value', axis=1)
3 y = data_encoded['median_house_value']
4
5 #4. Split data into 80% training and 20% testing sets
6 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

1 #5. Standardize the data
2 scaler = StandardScaler()
3 X_train_scaled = scaler.fit_transform(X_train)
4 X_test_scaled = scaler.transform(X_test)

```

```

1 #6. Perform Linear Regression
2 model = LinearRegression()
3
4 # Train the model on the training data
5 model.fit(X_train_scaled, y_train)
6
7 # Predict the target variable on the test data
8 y_pred = model.predict(X_test_scaled)
9
10 # Calculate Root Mean Squared Error (RMSE)
11 rmse = mean_squared_error(y_test, y_pred, squared=False)
12 print(f"Root Mean Squared Error (RMSE): {rmse}")

```

Root Mean Squared Error (RMSE): 70031.41991955665

```

1 #7. Perform Decision Tree Regression model
2 tree_reg_model = DecisionTreeRegressor()
3
4 # Fit the model on the training data
5 tree_reg_model.fit(X_train_scaled, y_train)
6
7 # Predict output for the test dataset
8 y_pred_tree = tree_reg_model.predict(X_test_scaled)
9
10 # Calculate Root Mean Squared Error (RMSE) for Decision Tree Regression
11 rmse_tree = mean_squared_error(y_test, y_pred_tree, squared=False)

```

```
12
13 print(f"RMSE (root mean squared error) from Random Forest Regression: {rmse_tree}")

RMSE (root mean squared error) from Random Forest Regression: 69228.49061115789
```

```
1 #8. Perform Random Forest Regression model
2 forest_reg_model = RandomForestRegressor()
3
4 # Fit the model on the training data
5 forest_reg_model.fit(X_train_scaled, y_train)
6
7 # Predict output for the test dataset
8 y_pred_forest = forest_reg_model.predict(X_test_scaled)
9
10 # Calculate Root Mean Squared Error (RMSE) for Random Forest Regression
11 rmse_forest = mean_squared_error(y_test, y_pred_forest, squared=False)
12
13 print(f"Random Forest Regression RMSE: {rmse_forest}")
```

Random Forest Regression RMSE: 48884.92495518843

```
1 #9. Plot the fitted model for test data
2 # Extract just the median_income column
3 X_train_median_income = X_train_scaled[:, X_train.columns.get_loc('median_income')]
4 X_test_median_income = X_test_scaled[:, X_test.columns.get_loc('median_income')]
5
6 # Reshape the data to fit the model
7 X_train_median_income = X_train_median_income.reshape(-1, 1)
8 X_test_median_income = X_test_median_income.reshape(-1, 1)
9
10 # Initialize Linear Regression model
11 linear_reg_single_var = LinearRegression()
12
13 # Fit the model on the training data
14 linear_reg_single_var.fit(X_train_median_income, y_train)
15
16 # Predict output for the test dataset
17 y_pred_single_var = linear_reg_single_var.predict(X_test_median_income)
18
19 # Plot the fitted model for training data
20 plt.figure(figsize=(10, 6))
21 plt.scatter(X_train_median_income, y_train, color='blue', label='Training Data')
22 plt.plot(X_train_median_income, linear_reg_single_var.predict(X_train_median_income), color='red', linewidth=3, label='Fitted Model')
23 plt.xlabel('Median Income')
24 plt.ylabel('Median House Value')
25 plt.title('Fitted Model - Training Data')
26 plt.legend()
27 plt.show()
```

```
28
29 # Plot the fitted model for test data
30 plt.figure(figsize=(10, 6))
31 plt.scatter(X_test_median_income, y_test, color='blue', label='Test Data')
32 plt.plot(X_test_median_income, y_pred_single_var, color='red', linewidth=3, label='Fitted Model')
33 plt.xlabel('Median Income')
34 plt.ylabel('Median House Value')
35 plt.title('Fitted Model - Test Data')
36 plt.legend()
37 plt.show()
```

1

